

Unobservable Messaging with MessageVortex

Martin Gwerder

Abstract—In this paper, we introduce an unobservable message anonymization protocol, named MessageVortex. It bases on the zero trust principle and a distributed peer-to-peer like architecture and avoids central aspects such as fixed infrastructures within a global network.

It scores over existing work by blending its traffic into suitable existing transport protocols, thus making it next to impossible to block it without significantly affecting regular users of the transport medium. No additional protocol-specific infrastructure is required in public networks and allows a sender to control all aspects of a message such as the degree of anonymity, timing, and redundancy of the message transport without disclosing any of these details to the routing or transporting nodes. The most recent RFC-draft describes the protocol in detail. This draft contains all the necessary information to build protocol nodes. More information and a prototype implementation are available under <https://messagevortex.net/>.

Index Terms—Computer Society, IEEE, IEEEtran, journal, L^AT_EX, paper, template.

1 INTRODUCTION

ALMON BROWN STROWGER was the owner of a funeral parlor in St. Petersburg. He filed a patent on March 10th, 1891 for an “Automatic Telephone Exchange” [1]. This patent built the base for modern automated telephone systems. According to several sources, he was annoyed by the fact that the local telephone operator was married to another undertaker. She diverted potential customers of Mr. Strowger to her husband instead, which caused Almon B. Strowger to lose business. In 1922, this telephone dialing system, which is nowadays called pulse dialing became the standard dialing technology for more than 70 years until tone dialing replaced it.

This dialing technology enabled automatic routing for voice and text messages (e.g., telex) up until today and is one of the foundations for our current routed networks. These networks build the base of our communication-based Society these days and allow us to connect quickly with any person or company of our wish. We use these networks today as communication meaning for all purposes, and most of the people spend minimal thoughts on the possible consequences arising if someone puts hands on this communication.

Collected data may be used to judge upon our intentions and thus is not only confidential if we have something to hide. This problem has dramatically increased in the last years as big companies and countries started to collect all kinds of data and created the means to process them. Such a judgment allows, supposedly, to classify people and their intentions. This is not limited to what they are doing but as well, on what they did and what they might do. Numerous events in the present and past show that multiple actors, some of which are state-sponsored, collected data on a broad base within the Internet. Whether this is a problem or not may be disputable. Undisputed is that such data requires careful handling, and accusations should then base

on solid facts. Unacceptable seems the use of “guesses” or “extrapolations” in most of the cases.

To show that this may happen even under complete democratic control, we may refer to events such as the “secret files scandal” (or “Fichenskandal”) in Switzerland. In the years from 1900 to 1990 Swiss government collected 900’000 files in a secret archive (covering roughly 10% of the natural and juristic entities within Switzerland at that time). The Swiss Federal Archives (<https://www.bar.admin.ch>) documents more about the Fichenskandal.

Whistleblower Edward Snowden leaked a vast amount of documents. These documents suggested that such attacks on privacy are commonly made on a global scale. The documents leaked in 2009 and a significant number of journalists from multiple countries screened them (e.g., NRC, the guardian, or Forbes). According to these documents, NSA infiltrated more than 50k computer networks with malware to collect classified or personal information. They furthermore infiltrated Telecom-Operators (according to the same documents mainly executed by British GCHQ) such as Belgacom to collect data and targeted high member of governments even in associated states (such as the mobile phone number of Germany’s president and high ranking members of economy and finance).

The National Security Agency (NSA) collected network data with a program called XKeyscore. Several proofs and in-depth information have been found within the documents leaked by Snowden (e.g., [2]). According to these papers XKeyscore spanned (in 2008) ≈ 150 sites with 700 Servers collecting emails, web traffic, and chat messages.

This list of events shows that big players are collecting and storing vast amounts of data for immediate or possible future analysis. The list of events also shows that the use of this data has in the past been at least partially questionable. As a part of possible countermeasures, this work analyses the possibility of using state-of-the-art technology to minimize the information footprint of a person on the Internet.

We leave a large information footprint in our daily communication. On a regular email, we disclose everything in a “postcard” to any entity on its way. Even when encrypting

• M. Gwerder was with the Institute of Mobile and Distributed Systems of the University of Applied Sciences of Northwestern Switzerland. E-mail: see <http://www.messagevortex.net>

Manuscript received June 1, 2019; revised July 20, 2019.

a message perfectly with today's technology (S/MIME [3] or PGP [4]) it still leaves at least the originating and the receiving entity disclosed, or we rely on the promises of a third party provider which offers a proprietary solution. Even in those cases, we leak pieces of information such as "message subject", "frequency of exchanged messages", "size of messages", or "client being used". A suitable anonymity protocol has, therefore, not only a message to hide but additional attributes as well. It includes leaving the message itself aside, all metadata, and all the traffic flows. Furthermore, a protocol to unlink and anonymize messages should not rely on the trust of infrastructure other than the infrastructure under control of the sending or receiving entity. Trust in any third party might be misleading in terms of security of the protocol.

Central infrastructure is bound to be of particular interest to anyone gathering data. It may furthermore allow manipulating the system or the data or the data flow. So, avoiding a central infrastructure is a good thing.

Leaving no information trail when sending information from one person to another is hard to achieve. Most messaging systems disclose at least the peer partners when sending messages. Metadata such as starting and endpoints, frequency, or message size are leaked in all standard systems even when encrypting messages.

Allowing an entity to collect data may affect senders and recipients of any information. Collection of vast amounts of data allows a potent adversary to build a profile of a person. Unlike in the past, the availability of this kind of information has been risen to a never known extend with the Internet.

An entity in possession of such Profiles may use them for many purposes. These include service adoption, directed advertising, or classification of citizens. The examples given above show that the effects of this data is not limited to the Internet but reaches us effectively in the real world.

The main problem of this data is that it may be collected over a considerable amount of time and evaluated at any time. It even happened that standard practices of the time are judged differently upon later. Persons may then be judged retrospectively upon these types of practice. This questionable type of judgment is visible in the tax avoidance discussion.

People must be able to control their data footprint. Not providing these means does effectively allow any country or a bigger player to ban and control any number of persons within or outside the Internet.

In this work, a new protocol is designed to allow message transfer through existing communication channels. These messages are next to unobservable for any third party. This unobservability does not only cover the message itself but all metadata and flows associated with it. We called this protocol "MessageVortex" or in short just "Vortex". The protocol is designed in such a way so that it is capable of using a wide variety of transport protocols. It is even possible to switch protocols while the messages are transferred. This behavior allows media breaches (at least on a protocol level) and makes analysis even harder.

The new protocol allows secure communication without the need for trusting the underlying transport media. Furthermore, the usage of the protocol itself is possible without altering the immediate behavior of the transport layer. That

way it is possible to use the transport layers regular traffic to increase the noise in which information has to be searched.

1.1 About Unlinkability, Unobservability, Undetectability and Censorship resistance

From an academic point of view, achieving anonymity is relatively simple. All we need is a trusted party distributing the messages while making sure that no trace from the sender arrives at the recipient. Unlinkability, as defined in [5] is much harder to achieve. It requires that a specified attacker is unable to link a sender and recipients of a message. Even if a system provides unlinkable message transfer, it is most probable prone to denial of service and thus partial or full censorship. By introducing a global observer or infiltrating parts of the system, an attacker may gain insight into the messages transported by the system and thus leaking information.

So to be censorship resistant, a protocol requires many critical, unobvious properties. As outlined, it should be undetectable from the outside. From within the system, we need to provide a mean to make it very hard, or ideally impossible to follow message flows or identify participants.

MessageVortex is a protocol providing censorship resistance under ideal circumstances. It does this using a rigid design from bottom up to provide the required properties. While being a protocol on its own, it uses many standard protocols. Partly to provide user-friendliness, but mostly to hide within the regular network flows. As such, a protocol requires to be undetectable on the network. A protocol all alone may not be undetectable as each protocol sends data over a network. This data is detectable. A protocol sending undetectable data, a censorship-resistant protocol requires to be embedded undetectably in legit message flows or hide in side channels. Such embedding is usually done either by side channel transmissions or by employing steganography. Our protocol may use both. The steganography is, however, the preferred way as it implies no control over the transport infrastructure.

1.2 Adversary Model

For our adversary, we need to assume someone of the capabilities of a state-sponsored actor. State-sponsored actors may work these days with allies. Achieving dominance on a world scale is excluded from our model.

For our adversary model we assume the following properties:

- An adversary has elaborated technical know-how to attack any infrastructure.
- An adversary may have the capability to monitor traffic at many points in most networks within a jurisdiction.
- An adversary may have the capability to modify routing information within a jurisdiction freely.
- An adversary may have the possibility to freely modify even cryptographically weak secured data where a single or a limited number of entities grant proof of authenticity or privacy.
- An adversary may have the possibility to inject or modify any data on the network of a jurisdiction.

- An adversary may create own nodes of a network. He may furthermore monitor their behavior and data flow to the maximum possible extent.
- An adversary may force a limited number of other non-allied nodes to expose their data to him.
- An adversary may have similar access to resources as within its jurisdiction in a limited number of other jurisdictions.

We assume the following goals for an adversary:

- An adversary may want to disrupt non-authorized communication.
- An adversary may want to read any information passing throughout the Internet.
- An adversary may want to build and conserve data about individuals or groups of individuals of their life. This data includes all activities in any part of their life regardless of their apparent fitness for any specific purpose or their correctness.

This adversary is much more potent than the supposed XKeyscore application. The Snowden papers suggest that this type of adversary is generally realistic and may exist.

1.3 Notation

The theory in this document is heavily based on symmetric encryption, asymmetric encryption, and hashing. To use an uniformed notation we use $E^{K_a}(M)$ (where a is an index to distinguish multiple keys) for an encrypting function with a key K_a . This results in M^{K_a} for the encrypted message. If we are reflecting a tuple of information, we write in boldface. To express a concatenated set of information, we use angular brackets $\langle \text{normalAddress}, \text{vortexAddress} \rangle$.

For a symmetric encryption of a message M with a key K_a resulting in M^{K_a} where a is an index to distinguish different keys. Decryption uses therefore $D^{K_a}(M^{K_a}) = M$.

As notation for asymmetric encryption we use $E^{K_a^1}(M)$ where as K_a^{-1} is the private key and K_a^1 is the public key of a key pair K_a^p . The asymmetric decryption is noted as $D^{K_a^{-1}}(M)$.

For hashing, we do use $H(M)$ if unsalted and H^{S_a} if using a salted hash with salt S_a . The generated hash is shown as H_M if unsalted and $H_M^{S_a}$ if salted.

If we want to express what details contained in a tuple we use the the notation $M\langle t, \text{MURB}, \text{serial} \rangle$ respectively if encrypted $M^{K_a}\langle t, \text{MURB}, \text{serial} \rangle$.

$$\begin{aligned}
 \text{asymmetric: } & E^{K_a^{-1}}(M) &= M^{K_a^{-1}} \\
 & D^{K_a^1}(E^{K_a^{-1}}(M)) &= M \\
 & D^{K_a^{-1}}(E^{K_a^1}(M)) &= M \\
 \text{symmetric: } & E^{K_a}(M) &= M^{K_a} \\
 & D^{K_a}(E^{K_a}(M)) &= M \\
 \text{hashing (unsalted): } & H(M) &= H_M \\
 \text{hashing (salted): } & H^{S_a}(M) &= H_M^{S_a}
 \end{aligned}$$

In general subscripts denote selectors to differentiate values of the same type and superscript denote relevant parameters to operations expressed. The subscripted and superscripted information may be omitted where not needed.

We refer to the components of a Vortex Message as follows:

$$\begin{aligned}
 \text{Prefix component: } & \text{PREFIX} &= D^{K_a^1}(\mathbf{P}^{K_a^{-1}}) = D(\mathbf{P}) \\
 \text{Header component: } & \text{HEAD} &= D^{K_a^1}(\mathbf{H}^{K_a^{-1}}) = D(\mathbf{H}) \\
 \text{Route component: } & \text{ROUTE} &= D^{K_a^1}(\mathbf{R}^{K_a^{-1}}) = D(\mathbf{R})
 \end{aligned}$$

In general a decrypted Block is written as capitalized multi character boldface. An encrypted Block is written as capitalized single character boldface.

To specify a random byte sequence of n octets generated by an PRNG of type t and seeded with s we use $R_t(s, n)$

1.4 Existing work

1.4.1 Definition of Censorship, Anonymity, Unlinkability and Censorship Resistance

As a definition for censorship, we take

Censorship: the cyclical suppression, banning, expurgation, or editing by an individual, institution, group or government that enforce or influence its decision against members of the public – of any written or pictorial materials which that individual, institution, group or government deems obscene and “utterly without redeeming social value,” as determined by “contemporary community standards.”

Which is attributed to Chuck Stone Professor at the School of Journalism and Mass Communication, University of North Carolina. Please note that “Self Censorship” (not expressing something in fear of consequences) is a form of censorship according to this definition too.

In our more technical context, we reduce the definition to

Censorship: A systematic suppression, modification, or banning of data in a network by either removal, or modification of the data, or systematic influencing of entities involved in the processing (e.g., by creating, routing, storing or reading) of this data.

This simplified definition narrows down the location to the Internet as it is the only relevant location to us. Furthermore, it limits the definition to the maximum reach within that system.

A censorship resistant system is a system which allows the entities of the system and the data itself to be unaffected from censorship. Please note that this does not deny the presence of censorship per se. It still exists outside the system. However, it has some consequences for the system itself.

- The system must be either undetectable or out of reach for an entity censoring.
- The possibility of identifying a protocol or data allows a censoring entity to suppress the use of the protocol itself.
- The entities involved in a system must be untraceable.

Traceable entities would result in a mean of suppressing real-world entities participating in the system.

For all other anonymity relevant terms, we refer to [5].. It defines anonymity as

Anonymity of a subject means that the subject is not identifiable within a set of subjects, the anonymity set.¹

and

Anonymity of a subject from an attacker's perspective means that the attacker cannot sufficiently identify the subject within a set of subjects, the anonymity set.¹

It defines unlinkability as

Unlinkability of two or more items of interest (IOIs, e.g., subjects, messages, actions, ...) from an attacker's perspective means that within the system (comprising these and possibly other items), the attacker cannot sufficiently distinguish whether these IOIs are related or not.¹

And undetectability as

Undetectability of an item of interest (IOI) from an attacker's perspective means that the attacker cannot sufficiently distinguish whether it exists or not.

For our work, We define the anonymity set as the set of all possible subjects within a supposed message. The anonymity of a subject towards an observing third party is a crucial factor as it relates directly to our adversary model.

The set of terms is further broadened with k -Anonymity as defined in [6].

1.5 Existing work regarding Anonymity Protocols and Censorship Resistance

We were unable to identify a protocol withstanding the strong definition of our adversary. We have, however, found many protocols dealing with anonymity and very few dealing censorship circumvention. We tried to learn from these approaches and tried to integrate useful approaches. It has to be said, however, that very little of the mentioned protocols here had ever experienced broad adoption. Even more, most of them were never implemented or challenged.

In terms of censorship circumvention, we found that only little work has been done in academia. Several technical ways have been explored to circumvent censorship. All of them seem to boil down to the following main ideas:

- Hide data
The most common approaches we have found were either mimicking protocols (as in [7]), use protocols as payload transports ([8]), or employ steganography (as in [9]) or comparable technologies as side channels or .
- Copy or distribute data to a vast amount of places in order to improve the lifespan of data
This has been done by systems like [10], or WikiLeaks.
- Outcurve censorship measurements
Censorship measurements, especially regarding the

Internet censorship of China, have been analyzed in depth under technological, sociological and economical aspects ([11], [12], or [13] to name a few)

For Anonymization the ideas seem mainly to concentrate around onionizing (ToR [14], SOR [15], DUO-Onions and Hydra-Onions [16]), DC networks (DC-Nets [17], Tarzan [18], GAS [8]), mixing (Babel [19], MorphMix [20], Mixminion [21], Salsa [22]) and DHT tables (Bifrost [23], BitBlender [24]).

As we use alien transport protocols instead of an own protocol, we decided to go for a mixing approach. This approach minimizes the number of messages. Unlike other protocols, we do not rely the logic of mixing on the routing node but entirely on the RBB.

2 THE MESSAGEVORTEX PROTOCOL

In this section, we introduce a new consistent, transport independent model for representing the different protocols used by MessageVortex. The focus of the description lies in academic concepts. For more technical information specifying the protocol as implemented, refer to [25].

Generally, the MessageVortex System consists only of nodes, whereas a node may be any system always connected to the Internet. This applies to any device regardless of NAT or similar technologies which usually oppose problems.

2.1 General Design

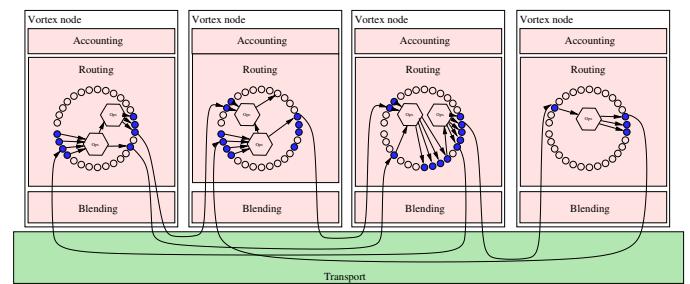


Fig. 1. A rough protocol outline of the MessageVortex protocol

The protocol handles messages which are passed by a transport protocol from node to node. The instructions how and when to pass a message is generated by a "routing block builder" (RBB). The RBB defines the path of the message, the type of hiding (blending) in the transport protocol, and the operations applied to each part of the message in each node. To avoid collision of operations, an RBB has "ephemeral identities" on each node with an assigned workspace. Ephemeral identities are short term identities containing a workspace and message quotas for a limited time. Ephemeral identities of a node are unrelated to each other. In the workspaces attached to the ephemeral identities, messages may be assembled, transformed, or decomposed with the operations above. Results are sent to other nodes. Due to the nature of the operations, all messages passed on may be decoy traffic or "real message parts" (We proof that claim in section 3.1). A node knows that a message is destined for it when a message results in data in specific parts of the identities workspace.

1. footnotes omitted in quote

As operations, we use some general capabilities such as splitting a message into two payload blocks and merging them again. Another type of operation is encryption and decryption of payload blocks. The third and most important type of operation is a redundancy operation. This operation uses a Reed-Solomon [26] function to add redundancy information to the data while obfuscating its content. This function has previously been proposed mainly for information sharing systems (e.g., [27]).

A routing block may be used once or multiple times if flagged accordingly. Repeating a routing block allows a sender to use a routing block as an anonymous endpoint address. It is essential to understand that reusing a routing block does have downsides in terms of privacy. Reusing a routing block does typically create the same pattern on the network assuming the same workspace layout. While the timing might vary the number of messages and the sequence of messages remains the same. For a full list of weaknesses when reusing routing blocks, see 3.2.

We divide our protocol into four different layers, whereas only three are specific to the MessageVortex protocol. The lowest layer is the transport layer. As expressed earlier, dedicated protocols are easy to censor. Therefore we build our protocol on top of other suitable transport protocols.

The other Three layers are Vortex-specific and do not require any infrastructure on the Internet. We elaborate further on these layers in section 2.3 and section 2.4.

For our first tests, we used a custom transport layer, allowing us to monitor all traffic quickly, and build structures in a very flexible way. This transport layer works locally with a minimum amount of work for setup and deployment. It furthermore works across multiple hosts in a broadcast domain. The API may be used to support almost any kind of transport layer. After that, we focused on the protocols identified as suitable as transport protocols:

- SMTP
- XMPP

For the prototype, we have implemented an SMTP transport agent and the respective blending layer.

The blending layer is taking care of multiple problems:

- Translating the message into the transport format
This translation includes jobs such as embedding a message as encoded text, as a binary attachment or hide it within a message using steganography.
- Extract incoming messages from the transport protocol
Identify incoming messages containing a possible block and extract it from the message.
- Do housekeeping on the storage layer
Access protocols such as POP and IMAP require message deletion after processing to stay below the sizing quotas of an account.

We define the blending layer to work as follows when receiving messages:

- 1) Log arrival time (in UTC) on the transport layer.
- 2) Extract possible blocks.
- 3) Apply decryption on a suspected header block.

- 4) Validate the header block using the accounting layer.
- 5) Process header requests (if any)
- 6) Extract and decrypt subsequent blocks.
- 7) Pass extracted blocks and information to the routing layer.

We define the blending layer to work as follows for sending messages:

- 1) Assemble message as passed on by the routing layer.
- 2) Using the blending method specified in the routing block build an empty message.
- 3) Create a message body content.
- 4) Send the message to the appropriate recipient using the transport layer protocol.

The routing layer receives the message blocks in a decrypted and authorized form from the blending layer. The routing layer then assembles all information of an identity and makes executes the accepted operations using the available data.

Tasks of the transport layer are:

- Build structure representing the block building and the appropriate block IDs.
- Schedule all Routing blocks for processing in a priority queue.
- Authorize all routing blocks ready for processing with the calculated block sizes.
- Process blocks.
- Send prepared building blocks to the Blending layer.

The only exclusions to these capabilities are made to the sending and the receiving node of a message. We always that these two nodes are not compromised observed.

2.1.1 Identities and Keys

During the life of a message, it requires several keys. In the following section, we emphasize on the type, the usage, and their specialties.

2.1.2 Peer key

The peer key is a message specific symmetrical key known to two adjacent routing nodes. The sending node generates it and encrypts it with the public key of the receiving nodes identity key.

2.1.3 Sender Key

The sender key is a symmetric key protecting the routing block and ephemeral identity information of the message. It is prepended to the header section and protected by the identity key of the processing node.

2.1.4 Identity Key

The Identity key is a static, asymmetric keypair existing on a per user base used to sign messages or encrypt symmetric keys. We refer here as a user any peer participating in the message stream. Every user participating in the Vortex network requires at least one key pair. Depending on the use-case (e.g., unlinkable signatures or scalable security), a user may use multiple key pairs at the same time.

2.1.5 Ephemeral Key

An ephemeral identity key identifies the sender to a processing host. It is asymmetric, short-lived (typically hours or days) and is created upon request of an RBB. It does not link to another identity (unlinkability). The purpose of an ephemeral key is to identify and protect the workspace on a node. This workspace is used for processing and storing VortexMessages and to process accounting information.

To create an ephemeral key, an RBB sends an empty VortexMessage to a node only containing a header request for creation of a new identity. The routing node may accept this request or reply with a request for proof of work. As a proof of work, the RBB has to send an identity block containing a specific pattern linked to the requested entity and the hash of the identity block has to start with a specified bit sequence. The length of the requested hash bit sequence reflects the complexity of the puzzle. The granting routing node may freely decide upon the validity and the complexity of the puzzle.

2.2 Messages and VortexMessages

A message is passed in full or only partially from routing node to another routing node. We refer as “message” to the full message as entered by the sender and received by the receiver. A “VortexMessage” is used for the message which is passed from routing node to routing node and contains the message, parts of the message or just decoy content. We show in figure 2 a simplified view on a VortexMessage.

A VortexMessage is passed from one router to another and is embedded as binary data in the transfer protocol. Every router may decide for himself on the support of algorithms and embedding mechanisms.

The block structure of a Vortex message is as follows:

- Encrypted peer key k_{peer_n}
Contains symmetrical key for decryption of follow up header information and payload blocks. The key is encrypted with the receiving host's public key. The key is known to the RBB, the message sending node, and the message receiving node. No other node knows the key.
- padding
This padding makes sure that the message part encrypted with the peer key does look different for each payload even when reusing keys. This case is not recommended but unavoidable in the case of a reused routing block.
- header block encrypted with k_{sender_n} and signed with the identities private key
 - Identity and “proof of work” information
This information contains the public key of the identity. It serves as an identifier for the workspace and is authenticates by the signature.
 - Replay protection information
This information allows a node to identify replayed Messages even if the payload of the content has been modified.
 - Forward secret information
This information allows a node to identify

VortexMessages where tampering occurred by recombining blocks of multiple messages.

- (optionally) Proof-of-work information
This information allows a sending node to fulfill a proof of work requirement raised due to a previously sent request.
- Routing blocks (encrypted with sender key)
 - Next hop timing instructions
This specifies by when building instructions should be carried out relative to the moment of reception of the VortexMessage. It is specified as a time range. There may be multiple timing instructions. Each of the instructions refers to precisely one routing block and one header block.
 - Next hop routing blocks
These routing blocks are placed into the VortexMessage created according to the build instructions and are already encrypted with $k_{sender_{n+1}}$ and thus not readable to the current node.
 - Next hop header (encrypted with $k_{sender_{n+1}}$)
These header blocks are placed into the VortexMessage created according to the build instructions and are already encrypted with $k_{sender_{n+1}}$ and thus not readable to the current node.
 - Message build instructions.
These instructions form the core for the workspace and contain all instructions and the information which payload blocks should be included in each of the messages.
 - Next hop peer key $k_{peer_{n+1}}$.
 - Next hop blending instructions.
These contain the information about what transport to use, what blending to use, and the address of the next router node.
- Encrypted payload blocks (encrypted with peer key)
 - Payload blocks

It is important to note that there are two symmetrical keys involved in encrypting and decrypting message headers. Having two keys is not a flaw in the protocol but necessary.

The first key of a VortexMessage is the peer key k_{sender_n} . This key is only accessible with the private key of the node receiving the message and is furthermore known by the RBB. It allows decryption of the routing blocks concerning the current node and the header information. The sender of a message block is therefore not able to tell if a message contains one or more routing blocks. It is important to note that no other node should have access to this information.

The second key is the peer key k_{peer_n} located in the encrypted header. The RBB chooses the key. This key protects the inner structure of the message. It makes it impossible for any node except the sending or the receiving peer node to detect the inner structure of the message. Without this key, any independent observer with knowledge about the blending capabilities of a receiving node may:

- easier identify the block structure.
This remains the case regardless of whether ASN.1 or

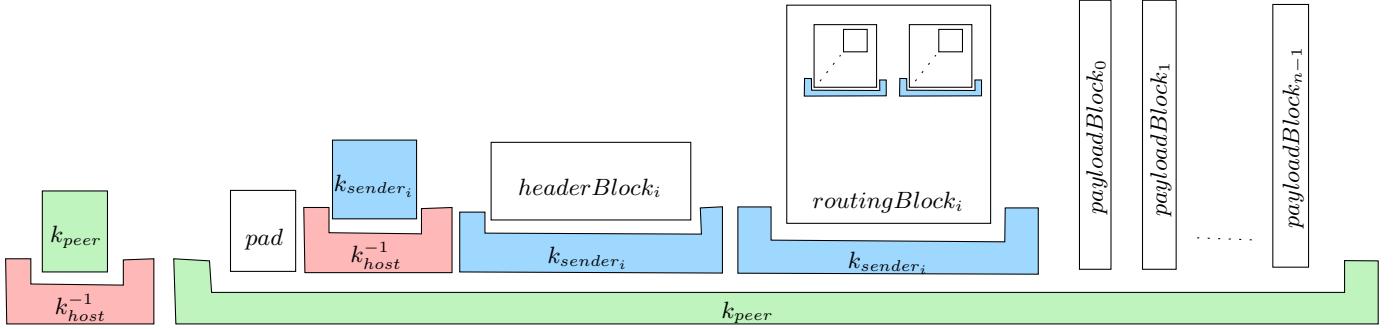


Fig. 2. Simplified message outline

length prefixed structures are used. If the structure of a vortex Message is identifiable, the messages may be logged or dropped by an adversary.

- Identify the routing block size.
The value of this information is only minimal as it only reflects the complexity of the remaining routing information indirectly.
- Identify the number of payload blocks and their respective sizes.

This is valuable information when following the traffic of a message.

2.3 Transport and Blending Layer

The tuple of Transport and Blending layer builds the main core functionality when it comes to the undetectability property. To avoid transport protocol misuse and unintentional exit nodes of the protocol, the RBB has no control over the transported content except for the hidden VortexMessage and how it is embedded. This rule loads the burden of sensible cleartext payload generation to the blending layer.

It is relatively easy to generate a credible cleartext message to pass an automated testing engine. This statement may be verified by looking at the effectiveness of today's junk mail filters. These filters have huge problems continuously adapting to the new types of unsolicited bulk emails (UBE).

Things do, however, drastically change if taking a human censor into account. A human censor is not only able to analyze the text and layout of a message. He is furthermore capable of judging on the stringency of a communication. He may deduce data such as relationship and type of writing. Then, he may detect anomalies within conversations and judge whether the communication pattern is more likely to be from a human or a chatbot.

2.3.1 Applied Steganography and the Dead Parrot

A human censor can take very complex information into account when it comes to analyses of message content. He is not only able to analyze a message for its content, but he may also see the message in the context of other messages. In [28] is expressed that it is easy for a human to determine decoy traffic as the content is easily identifiable as generated content. While this is true for the very general case, there is a possibility here to generate "human-like" data traffic to a certain extent. As an adversary may not assume that his messages are replied to, the problem does not boil down to a Turing test. It remains on the level of a "passive observer

Turing test", in this scenario the censor is only able to judge on the given messages instead of introducing his own questions, wordings, and verbal challenges. By enabling the potential nodes to choose their messages and the replies generated to them, we enabled them to choose very reduced types of communications. The chosen messages may even be identifiable as automated messages (e.g., messages of a monitoring system or messages of an SMS to email gateway).

The most straightforward approach would have been to give a routing block builder the means of controlling the decoy message content. While such a possibility would be easy, it would enable a routing block builder to use the node as a "exit node" from the system. Blackmailing messages could be sent through the system to a non-participating member and leak at the same time the presence of a routing node. To deny this possibility, we shifted the ability to the routing node.

The VortexMessage itself is binary, and as such, there are only limited possibilities to hide it within the transport protocol. We decided to use attachments or attachment-like structures. Within the attachments, we currently support two types of embedding: plain and steganographic embedding.

Plain embedding means that we insert a sequence of blocks into a standard message. This is typically done within files with a weaker structure and high entropy (such as an MP3 encoded file). While this is very hard to detect for a machine, it becomes immediately suspicious for a human censor. A human censor would detect the presence of a payload which does not make any sense.

For steganographic embedding, we decided to go for F5 [9]. It is a reasonably well-researched algorithm which attracted many researchers. The original F5 implementation had a detectable issue with artifacts [29] caused by the recompression of the image. This issue was caused only due to an issue in the reference implementation, and the researchers have provided a corrected reference implementation without the weakness. Like always, the type of embedding may be specified and replaced upon request.

2.4 Routing and Accounting

Each routing layer does the anonymization operations. These operations should make sure that no message may be identified unless destined for the current node. An incoming VortexMessage is first extracted by the routing layer and

authenticated through the accounting layer. For a message to be valid, the identity must be known (or include an identity creation request), the VortexMessage must not be replayed, and there must be an available quota for the message. If all criteria apply the VortexMessage content consisting of routing blocks, operations, and the payload is added to the workspace of the identity.

Each routing block of a workspace has an active time range measured in seconds since the arrival of its message on the transport layer. If a routing block becomes active, operations for building its payload blocks are processed, and new VortexMessages are assembled and passed to the blending layer for embedding and sending.

2.5 Header Requests

There are several header requests available to the protocol mainly supporting the message transfer. There are header requests for creating ephemeral identities on a node, and for raising a message or size quota. Additionally, there are some helper requests for querying quotas of an ephemeral identity and for querying the capabilities of a node. Additionally, a node may offer knowledge about other routing nodes. This request should, however, only discover nodes in jurisdictions, where the operation of such nodes is legal.

Some of the requests may be replied with the requirement of providing a “proof of work” to be fulfilled. A node is free to not reply to a request or reject it depending on its privacy needs and the current usage of resources.

2.6 The Core: Operations

We differentiate three types of operations:

- Splitting and merging of chunks
- Encryption and decryption of chunks
- Redundancy calculations carried out on chunks

The first two Operations do not provide a high level of unlinkability as they do allow analysis such as hotspot analysis and produce continuously inclining, steady or declining message sizes depending on the type of use. The third operation, however, adds a whole lot of new possibilities in conjunction with the other two.

2.6.1 Splitting and Merging

The splitPayload operation splits a payload block into two chunks of different or equal sizes. The parameters for this operation are:

- source payload block pb_1
- fraction f

A floating-point number which is describing the size of the first chunk. If the fraction is “1.0”, then the whole payload is transferred to the second target chunk

If $len(pb_1)$ expresses the size of a payload block called pb_1 in bytes, then the two resulting blocks of the SpitPayload Operation pb_2 and pb_3 have to follow the following rules:

$$split(f, pb_1) = \langle pb_1, pb_2 \rangle \quad (1)$$

$$pb_1.startsWith(pb_2) \quad (2)$$

$$pb_1.endsWith(pb_3) \quad (3)$$

$$len(pb_2) = \lfloor len(pb_1) \cdot f \rfloor \quad (4)$$

$$len(pb_1) = len(pb_2) + len(pb_3) \quad (5)$$

The mergePayload operation combines two payload blocks into one. The parameters for this operation are:

- first source payload block pb_1
- second source payload block pb_2

If $len(pb)$ expresses the size of a payload block called pb in bytes then resulting block of the MergePayload Operation pb_3 have to follow the following rules:

$$merge(pb_1, pb_2) = pb_3 \quad (6)$$

$$pb_3.startsWith(pb_1) \quad (7)$$

$$pb_3.endsWith(pb_2) \quad (8)$$

$$len(pb_3) = len(pb_1) + len(pb_2) \quad (9)$$

2.6.2 Encryption and Decryption

The encryptPayload operation encrypts a payload block pb_1 symmetrically resulting in a block pb_2 . The length of block pb_2 may vary according to mode and padding chosen. The parameters for this operation are:

- Source payload block pb_1
- Encryption specification $spec$
- Symmetric key k

The operation follows the following rules (please note section 1.3 for notation):

$$encrypt(pb_1, spec, k) = pb_2 \quad (10)$$

$$pb_2 = E_{spec}^{K_a}(pb_1) \quad (11)$$

$$len(pb_2) \geq len(pb_1) \quad (12)$$

The decryptPayload operation decrypts a payload block pb_1 symmetrically resulting in a block pb_2 . The length of block pb_2 may vary according to mode and padding chosen. The parameters for this operation are:

- Source payload block pb_1
- Decryption specification $spec$
- Symmetric key k

The operation follows the following rules (please note section 1.3 for notation):

$$decrypt(pb_1, spec, k) = pb_2 \quad (13)$$

$$pb_2 = D_{spec}^{K_a}(pb_1) \quad (14)$$

$$len(pb_2) \leq len(pb_1) \quad (15)$$

2.6.3 Redundancy Operations

These operations build the core of the mixing operations. The operation allows to add to a message redundancy information or to rebuild a block from a chosen set of information.

The operation itself is shown in Fig. 3.

It may be subdivided into the following operations:

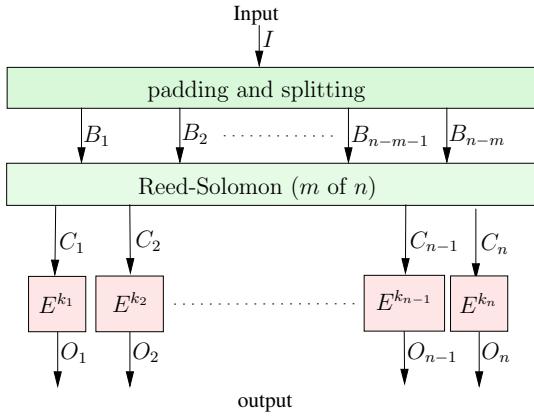


Fig. 3. Outline of the addRedundancy operation

- Pad the original message block in such a way, that all resulting blocks are a multiple of the block size of the encrypting cipher.
- Apply a Reed Solomon operation in a given GF space with a Vandermonde matrix.
- Encrypt all resulting blocks with unpadded, symmetrical encryption.

The padding is not a standard padding from encryption. The reason for this lies in the properties required in the padding. These properties were:

- The padding must not leak whether the rebuild cycle was successful or not.
- The padding should only reveal true message size if a message was successfully rebuilt.
- The anyone knowing the routing block content and the transmitted message must be able to predict any treated block including all padding bytes
- The padding must work with any size padding space

The padded value \mathbf{X} for a function $RS_{m \text{ of } n}$ and an encryption block size of k is calculated as follows:

$$i = \text{len}(\mathbf{M}) \quad (16)$$

$$e = k \cdot n \quad (17)$$

$$l = \left\lceil \frac{i + 4 + C2}{e} \right\rceil \cdot e \quad (18)$$

$$p = i + \left(C1 \cdot l \pmod{\left\lfloor \frac{2^{32} - i}{l} \right\rfloor \cdot l} \right) \quad (19)$$

$$\mathbf{X} = \langle p, \mathbf{M}, R_t(s, l - i - 4) \rangle \quad (20)$$

The remainder of the input block, up to length l , is padded with random data. The random padding data may be specified by RBB though a PRNG spec t and an initial seed value s . The message is padded up to size L . All resulting, encrypted blocks do not require any padding. This because the initial padding guarantees that all resulting blocks are dividable by the block size of the encrypting function. If not provided by an RBB, an additional parameter $C1$ is chosen as random positive integer and $C2 = 0$ by the node executing the operation.

To reverse a successful message recovery information the of a padded block \mathbf{X} , we calculate the original message size by extracting p and doing $\text{len}(\mathbf{M}) = p \pmod{\text{len}(\mathbf{X})}$.

This padding has many important advantages:

- The padding does not leak if the rebuilding of the original message was successful. Any value in the padding may reflect a valid value.
- Since we have a value $C2$, the statement that a message size is within $\text{len}(\mathbf{X}) < \text{size} < (\text{len}(\mathbf{X}) - k \cdot n)$ is no longer true and any value smaller $\text{len}(\mathbf{X}) - k \cdot n$ may be correct as well.
- An RBB may predict the exact binary image of the padded message when specifying $C1$, $C2$, and $R_t(s, .)$.

The Reed-Solomon operation is done with a Vandermonde matrix. Unlike in error correcting systems, we do not normalize the matrix so that the result of the first blocks is equivalent to the original message. Instead, the error correcting information is equally distributed over all resulting blocks adding further obfuscation. Since the entropy of the resulting blocks is lowered and may thus leak an estimate of how a resulting block may have been treated, we added the encryption step to equalize entropy again. The previously introduced padding guarantees that there is no further padding on block level required. Such padding could leak in case of decryption whether the block has been altered or not.

2.7 Usage of the Protocol

First, a sending node collects either a set of nodes and keys it wants to use and creates identities on these nodes using header requests. Then the sender creates a routing block containing all the routing instructions (hops and operations). Alternatively, a sender may use a premanufactured routing block for the specified target. This routing block is then concatenated to a message and passed to the locally running routing node. From there the message is routed as defined in the routing block. An example of such a route is shown in Fig. 1.

3 DISCUSSION OF THE RESULTS

We first focus on the protocol itself to show the strength and weaknesses of the protocol. After that, we focus on the dynamic part and see what type of data may be collected when considering not only the protocol but the whole message flow. We then present guidelines for different jurisdictional types.

We focus on an adversary in an environment, where the participation as a MessageVortex router, is considered a criminal act and highlight some additional constraints applying in such situations.

3.1 Static Protocol Analysis

A VortexMessage is not identifiable as the message is structureless on the outside. The VortexMessage itself follows the encrypted key without any structure. Therefore, we require the hosts private key to tell whether there is VortexMessage within a transport message or not.

The padding block PAD makes sure that, even if a routing block is reused, the VortexMessage structure is not the same. However, the preceding block with the key

remains the same. This block may serve as a fingerprint for an adversary to identify repeated MURBs.

Next, and one of the biggest problems we found is that a VortexNode is aware of its immediate peers. This flaw is because we do require a routable address for the transport protocol. Vortex nodes may thus discover their peers. It is, on the other side, not possible to use discovered peers. To use a peer, a transport address and a host key are required. A VortexNode may query this key, but there is no obligation to reply for the node asked for the key.

An active adversary may not create its routing blocks or header blocks and inject them due to the forward secret. He may, however, replace the peer key of a message. As this key is known to him, he gains no additional knowledge. Replacing the sender key block breaks the message. Replacing the header or routing block of the message with another header or routing block from the same ephemeral identity breaks the message unless the RBB reused the sender key and the forward secret. Finally, exchanging, omitting, or adding payload blocks renders the message inoperable, but does not generate additional knowledge. Replying the same or a modified block does not generate any pattern on the network as the replay protection stops propagating messages at the next node. Thus, a replayed block does not generate new knowledge to an observer.

All operations may apply to true message chunks as well as decoy traffic. As a node cannot tell if a traffic arriving is a decoy or true message content, it is unable to tell apart what outgoing traffic is a decoy. An encrypted block is of the same nature before and after encryption. As we do not know the blocks nature before, we are unable to tell the blocks nature after the encryption. The same argument applies to decryption, split, and merge operations.

Redundancy operations are alike. They, however, fulfill an additional purpose. A *addRedundancy* operation adds size to a message without differentiating between redundancy information and original payload. If the original block was a decoy, then all resulting blocks are decoys. If an originating block was message content, then all resulting blocks hold the same amount of data from the original block. So, this operation allows decoy traffic generation without enabling a generating node to identify the decoy traffic.

3.1.1 Endpoint Operation

Depending on the blending method, an adversary may identify single messages as long as they are detectable. Detectability depends on various factors, such as:

- Broken internal file structure (due to plain blending)
- Uncommon high entropy in a structureless file
- Unrelated message flow (see [28])
- Non-human behaviour on the transport layer (e.g., message traffic 24x7)

If an adversary identifies an endpoint successfully, then all peering endpoints of the same protocol may be identified as well by following the message flow. This does, however, not enable an adversary to inject messages as the host key is not known.

Assuming a global observer as an adversary and unencrypted traffic, he might discover the originating routing layer and thus identify it as Vortex node by following

traces of the transport layer. In most protocols, however, this address is spoofable and not a reliable source for the originating account.

3.1.2 Conclusions based on ephemeral identities

The knowledge a node may gain from ephemeral identities is minimal. The ephemeral identity is created by a node unknown to the receiver of the request. The only thing we know is what node was adjacent when creating the ephemeral identity. As the creation of an ephemeral identity is not linked to any other identity or ephemeral identity relationship between ephemeral identities on two nodes cannot be established. If two adjacent nodes cooperate when processing two linked ephemeral identities, no additional knowledge may be won. If two collaborating nodes have one or more non-collaborating nodes between them, they lose all linking knowledge due to the non-collaborating nodes.

3.1.3 Conclusions from Operations

Operations have been carefully crafted to leak as little information as possible. Being able to encrypt or decrypt a payload block does not leak any information. The data processed may be true message traffic or decoy as we do not know what the nature of the received message was. If an RBB avoids repeating patterns on nodes, it is not possible to link ephemeral identities of two non-adjacent nodes. Repeating patterns may arise, for example, if a block pb_1 is decrypted and re-encrypted on two nodes. In this case, both nodes may match the message as it contains the same content between the operations.

node f:

$$\begin{aligned} pb_2 &= D(pb_1) \\ pb_3 &= E^{K_t}(pb_2) \end{aligned}$$

node f+1:

node f+x:

$$pb_4 = D^{K_t}(pb_3)$$

In this example the patterns of pb_3 and $pb_4 = pb_2$ are two patterns repeating on non-adjacent nodes. The same conclusions are even more valid for splitting operations. These two operations should be regarded as helpers for the *addRedundancy* and *removeRedundancy* operations. These operations may be used to generate decoy traffic or to destroy data without knowledge of doing so of the processing node. If we process a function $addRedundancy_{2of3}$ any of the output blocks contains the input payload and any two of them may be used to recover the data. At the same time, an operation $removeRedundancy_{2of3}$ may be successful or not. The node is unable to differentiate between the two states. The padding applied and the unpadded encryption makes it impossible to judge upon success or fail of an operation.

3.2 Dynamic Protocol Analysis

A global observer is unable to analyze a message flow by timing or pattern of the exchanged packet even when being able to identify message vortex packages. Entry and exit nodes are indistinguishable even if having infiltrated significant portions of the network. Cooperation between adjacent nodes does not gain more information as all operations are minimized then to one combined workspace with all the operations. Linking of the message of two non-adjacent nodes is not possible as there are no linking attributes.

3.2.1 Bootstrapping of Addresses and Identities

Using the header requests an adversary may discover nodes over time. While it is not possible to screen traffic destined to such nodes, a global observer may identify peer partners of these nodes on the transport level.

3.2.2 Discovery of Peer Nodes

Besides attacking the message content, attacking the routing nodes is an option for an adversary in a jurisdiction where the operation of such a node is a criminal act. The presence of routing nodes is valuable information for a global observer narrowing down the information regarding data to be analyzed.

3.2.3 Findings based on adversary environment

In environments containing only global observers and no jurisdictional constraints regarding the technology, a VortexNode may disclose its presence. This means a VortexNode is not forced to cloak its presence. In such an environment, an RBB should choose the operations to be sensible, but extensive care is not required. Even if there is a node with a known owner of the node and a suspected message is received, the owner may credibly claim that the message in question was a decoy. No information obtained by any node involved in the routing of the message may proof anything else. Since a message may be split into any number of parts and related messages are only identifiable with a high degree of improbability even meta information such as the real size of the message, the sending time or the involved parties in the anonymity set are unknown. This statement is still valid if we consider an active adversary.

In environments where using a VortexNode is subject to criminal prosecution, much more care has to be applied. As all routing nodes know their immediate peer, we were only able to find two weak solutions to this problem. The first solution is only to use trusted nodes. If we can trust all routing nodes, no external observer may prove that the message flow is, in fact, MessageVortex traffic. An adversary node within such a system can learn the addresses of other nodes within the set. The RBB may reduce the set of uncovered nodes by applying communicating groups of nodes (communication cells) with defined gateways nodes between them. In such a scenario, only a cell and a possible adjacent cell may be discovered.

3.2.4 Issues When Reusing Routing Blocks

Reusing a routing block is required if the receiver is not known, and a continuous stream of messages is required.

Although it is possible to use multiple single route routing blocks (SURB) instead of one multi use routing block (MURB), it is costly. These costs arise due to the necessary calculation power to create identities. MURBs do have, however, significant drawbacks in terms of unlinkability and should be, therefore, avoided if possible.

A MURB creates a repeated pattern on the network in terms of messages. For a routing node, it is evident that the same tuple of communication partners is exchanging messages. The size of the VortexMessage allows in such a case an estimate of size in relation to the previous messages.

Furthermore, security is affected when using MURBs. A MURB may be replayed and allows thus to exhaust quotas of an ephemeral identity. To counter such exhaustion, the protocol introduces a maximum replay rate, but this is only weak protection.

3.3 The Missing Links

A lot still needs to be done. For this protocol to be of any use, a user-friendly implementation is required. The currently released implementation works as a prototype for academic research. It is, however, far beyond from being user-friendly. Too many choices are currently left to the user. A new implementation must provide excellent censorship resistance while providing easy to use recipes for message transfer.

For the traffic to be truly undetectable, chatbots must generate meaningful conversation between blending nodes. This conversation does not necessarily boil down to a Turing test. It is sufficient that two blending layers are capable of setting up communication, which is indistinguishable from a regular human or machine communication. As an adversary is typically not able to generate own traffic without exposing the probing activity and a blender is not required to such probes, an attacker is very limited.

Furthermore, some issues have been identified relating to updating nodes. A node should be able to request the software over VortexMessages as official sources for updates may be blocked.

Next is the problem of the routing node. The hardware of such a node should be protected with a small platform featuring deniable encryption and anti-forensic measures. We are currently investigating the possibility of creating such a cheap platform based on a RaspberryPi Zero.

Another exciting field of academic research is creating strategies for Routing block builders (RBB). We currently have a toolset of powerful operations, but academic researched strategies or guidelines for good routing blocks are missing.

4 CONCLUSION

Creating a protocol which is possibly censorship resistant, is already hard. The analysis showed that even when a protocol is crafted with great care, braking unobservability is far simpler than doing it right. At the same time, we were able to present a system which requires an unmatched amount of observation, infrastructure, and calculation power to be broken.

REFERENCES

- [1] A. B. Strowger, "Automatic telephone-exchange," United State Patent Office Patent 447918, 3 10, 1891.
- [2] NSA, Web and several newspapers, Jul. 2013, three slides have been redacted as they contained supposedly specific NSA operations. [Online]. Available: <https://www.theguardian.com/world/interactive/2013/jul/31/nsa-xkeyscore-program-full-presentation>
- [3] N. Freed and N. Borenstein, *RFC2045 Multipurpose Internet Mail Extensions; (MIME) Part One: Format of Internet Message Bodies*. IETF, 1996. [Online]. Available: <http://tools.ietf.org/pdf/rfc2045.pdf>
- [4] M. Elkins, *RFC2015 MIME Security with Pretty Good Privacy (PGP)*. IETF, 1996. [Online]. Available: <http://tools.ietf.org/pdf/rfc2015.pdf>
- [5] A. Pfitzmann and M. Hansen, "A terminology for talking about privacy by data minimization: Anonymity, unlinkability, undetectability, unobservability, pseudonymity, and identity management," http://dud.inf.tu-dresden.de/literatur/Anon_Terminology_v0.34.pdf, Aug. 2010, v0.34. [Online]. Available: http://dud.inf.tu-dresden.de/literatur/Anon_Terminology_v0.34.pdf
- [6] L. von Ahn, A. Bortz, and N. J. Hopper, "k-anonymous message transmission," in *Proceedings of the 10th ACM Conference on Computer and Communications Security (CCS 2003)*, V. Atluri and P. Liu, Eds. ACM Press, Oct. 2003, pp. 122–130. [Online]. Available: <http://www.abortz.com/papers/k-anon.pdf>
- [7] H. Mohajeri Moghaddam, "Skypemorph: Protocol obfuscation for censorship resistance," Master's thesis, University of Waterloo, 2013. [Online]. Available: https://uwspace.uwaterloo.ca/bitstream/handle/10012/7262/Mohajeri_Hooman.pdf?sequence=1
- [8] E. Athanasopoulos, M. Roussopoulos, K. G. Anagnostakis, and E. P. Markatos, "Gas: Overloading a file sharing network as an anonymizing system," in *Proceedings of Second International Workshop on Security, (IWSEC 2007)*, 2007. [Online]. Available: <http://dcs.ics.forth.gr/Activities/papers/gas%20iwsec07.pdf>
- [9] A. Westfeld, "F5 - a steganographic algorithm," *none*, vol. none, 2002. [Online]. Available: <http://www.ws.binghamton.edu/fridrich/research/f5.pdf>
- [10] I. Clarke, O. Sandberg, B. Wiley, and T. W. Hong, "Freenet: A distributed anonymous information storage and retrieval system," in *Proceedings of Designing Privacy Enhancing Technologies: Workshop on Design Issues in Anonymity and Unobservability*, Jul. 2000, pp. 46–66. [Online]. Available: <https://freenetproject.org/>
- [11] R. Ensaifi, D. Fifield, P. Winter, N. Feamster, N. Weaver, and V. Paxson, "Examining how the great firewall discovers hidden circumvention servers," 2015. [Online]. Available: <https://www.nymity.ch/papers/pdf/ensafi2015b-slides.pdf>
- [12] R. Clayton, S. J. Murdoch, and R. N. M. Watson, "Ignoring the great firewall of china," in *Privacy Enhancing Technologies*. Springer Berlin Heidelberg, 2006, pp. 20–35.
- [13] G. Lowe, P. Winters, and M. L. Marcus, "The great dns wall of china," *MS, New York University*, vol. 21, p. 1, 2007. [Online]. Available: <https://censorbib.nymity.ch/pdf/Lowe2007a.pdf>
- [14] R. Dingledine and N. Mathewson, "Tor protocol specification." [Online]. Available: <https://gitweb.torproject.org/torspec.git/tree/tor-spec.txt>
- [15] A. Egners, D. Gatzen, A. Panchenko, and U. Meyer, "Introducing sor: Ssh-based onion routing," in *Advanced Information Networking and Applications Workshops (WAINA), 2012 26th International Conference on*, IEEE, Mar. 2012, pp. 280–286. [Online]. Available: https://www.researchgate.net/profile/Andre_Egners/publication/237007773_Introducing_SOR_SSH-based_onion_routing/links/548805e90cf2ef34478ed724/Introducing-SOR-SSH-based-onion-routing.pdf
- [16] J. Iwanik, M. Klonowski, and M. Kutyłowski, "Duo-onions and hydra-onions-failure and adversary resistant onion protocols," in *Communications and Multimedia Security*. Springer, 2005, pp. 1–15. [Online]. Available: <http://dl.ifip.org/db/conf/cms/cms2004/IwanikKK04.pdf>
- [17] D. Chaum, "The dining cryptographers problem: Unconditional sender and recipient untraceability," *Journal of Cryptology*, vol. 1, pp. 65–75, 1988. [Online]. Available: <http://www.cs.ucsb.edu/~ravenben/classes/595n-s07/papers/dcnet-jcrypt88.pdf>
- [18] M. J. Freedman and R. Morris, "Tarzan: A peer-to-peer anonymizing network layer," in *Proceedings of the 9th ACM Conference on Computer and Communications Security (CCS 2002)*, Nov. 2002. [Online]. Available: <http://pdos.lcs.mit.edu/tarzan/docs/tarzan-ccs02.pdf>
- [19] C. Gülcü and G. Tsudik, "Mixing E-mail with Babel," in *Proceedings of the Network and Distributed Security Symposium - NDSS '96*. IEEE, Feb. 1996, pp. 2–16. [Online]. Available: <http://citeseer.nj.nec.com/2254.html>
- [20] M. Rennhard and B. Plattner, "Introducing morphmix: Peer-to-peer based anonymous internet usage with collusion detection," in *Proceedings of the Workshop on Privacy in the Electronic Society (WPES 2002)*, Nov. 2002. [Online]. Available: <http://cecid.sourceforge.net/morphmix.pdf>
- [21] G. Danezis, R. Dingledine, and N. Mathewson, "Mixminion: Design of a type iii anonymous remailer protocol," in *Proceedings of the 2003 IEEE Symposium on Security and Privacy*, May 2003, pp. 2–15. [Online]. Available: <http://mixminion.net/minion-design.pdf>
- [22] A. Nambiar and M. Wright, "Salsa: A structured approach to large-scale anonymity," in *Proceedings of CCS 2006*, Nov. 2006. [Online]. Available: <http://ranger.uta.edu/~mwright/papers/salsa-ccs06.pdf>
- [23] M. Kondo, S. Saito, K. Ishiguro, H. Tanaka, and H. Matsuo, "Bifrost : A novel anonymous communication system with DHT," in *2009 International Conference on Parallel and Distributed Computing, Applications and Technologies*. IEEE, dec 2009.
- [24] K. Bauer, D. McCoy, D. Grunwald, and D. Sicker, "Bitblender: Light-weight anonymity for bittorrent," 2008.
- [25] M. Gwerder, "Messagevortex protocol," IETF, 2018. [Online]. Available: <https://messagevortex.net/documentation.html>
- [26] I. S. Reed and G. Solomon, "Polynomial codes over certain finite fields," *Journal of the society for industrial and applied mathematics*, vol. 8, no. 2, pp. 300–304, Jun. 1960. [Online]. Available: <https://faculty.math.illinois.edu/~duursma/CT/RS-1960.pdf>
- [27] R. J. McEliece and D. V. Sarwate, "On sharing secrets and reed-solomon codes," *Communications of the ACM*, vol. 24, no. 9, pp. 583–584, 1981. [Online]. Available: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.80.2829&rep=rep1&type=pdf>
- [28] A. Houmansadr, C. Brubaker, and V. Shmatikov, "The parrot is dead: Observing unobservable network communications," in *Proceedings of the 2013 IEEE Symposium on Security and Privacy*, May 2013. [Online]. Available: <http://www.cs.utexas.edu/~amir/papers/parrot.pdf>
- [29] J. A. Briffa, H. G. Schaathun, and A. W. A. Wahab, "Has f5 really been broken." [Online]. Available: <https://pdfs.semanticscholar.org/4d6c/d9d7e3a419ea74a4a363a36fcc674e89ecc7.pdf>

Martin Gwerder

Martin Gwerder was born 20. July 1972 in Glarus, Switzerland. He is currently a doctoral Student at the University of Basel. After having concluded his studies at the polytechnic at Brugg in 1997, he did a postgraduate study as a master of business and engineering. Following that, he changed to the university track doing an MSc in Informatics at FernUniversität in Hagen. While doing this he constantly broadened his horizon by working for industry, banking and government as engineer and architect in security related positions. He currently holds a lecturer position for cloud and security at the University of Applied Sciences Northwestern Switzerland. His main expertise lays in the field of networking related problems dealing with data protection, distribution, confidentiality and anonymity.

